

Diotima: Compliance-Grade AI Infrastructure for Formative Assessment in Education

NOVAL CONSULTING

AI Consulting and EU AI Act Advisory Practice, Ireland. Contact: info@novalconsultancy.com

ABSTRACT

The EU AI Act's classification of student assessment systems as high-risk under Annex III creates both a compliance obligation and a need to answer a strategic design question: what does an educational AI platform look like when it is built for regulatory alignment from first principles, rather than retrofitted to it? This paper describes Diotima, a teacher-centred formative assessment platform that treats Annex III obligations not as constraints but as architectural foundations. By grounding all AI-generated content in approved curriculum materials, scoping high-risk components narrowly, enforcing human oversight as a structural requirement of the system rather than a policy statement, and embedding post-market monitoring as an operational discipline, Diotima demonstrates that high-risk AI in education can be simultaneously pedagogically sound, technically robust, and legally compliant. The paper argues that compliance-by-design is not a cost of participation in regulated education markets—it is the only credible path to adoption at scale.

Keywords EU AI Act Annex III Formative assessment Educational technology Human oversight Compliance-by-design

1. INTRODUCTION

In proposing a new framework for AI governance in education, the EU AI Act draws a conclusion that practitioners have long recognised implicitly: assessment is not a neutral activity. EUROPEAN PARLIAMENT (2024, Art. 6) classifies AI systems used to assess students in formal educational settings as high-risk under Annex III, Section 3(a). This classification reflects the reality that assessment shapes student progression and opportunity, that biased or incorrect automated feedback can cause material harm to learners, and that ungoverned automated evaluation can progressively displace the professional judgment on which good teaching depends.

Yet the full implications of this regulatory position have not yet been fully absorbed by the market. Evidence suggests that most educational AI products treat compliance as a downstream concern—a documentation task to be completed once the product is built, rather than a constraint that should shape every architectural decision from the outset. A governance assessment of AI-based assessment systems across Italian universities found that only 13% of institutions had adopted AI policies by September 2025, and that existing documentation prioritised aspirational statements over operational mechanisms (FERRARIO ET AL., 2025). At the international level, an OECD review of 18 countries found that none had issued specific regulation on the use of generative AI in education as of early 2024 (OECD, 2023).

The purpose of this paper is to describe a different approach. Diotima, a teacher-centred formative assessment platform, serves as a case study for a set of design principles that are more broadly applicable to any high-risk educational AI system. Designed from first principles as a compliance-grade system—one that treats Annex III obligations not as a constraint on product development but as the specification for it—Diotima demonstrates that the architecture required for regulatory alignment is also the architecture required for pedagogically serious, professionally empowering educational

AI. The paper describes Diotima's approach to human oversight and explainability, its model governance framework, and its data governance posture, and examines the implications of compliance-by-design for adoption in regulated education markets.

The paper proceeds as follows. Section 2 establishes the regulatory context and identifies three foundational challenges that any serious educational AI platform must address. Section 3 describes Diotima's core design proposition and how it responds to each of those challenges. Section 4 provides a system overview, illustrated by the lifecycle diagram in Figure 1 and the governance mapping in Table 1. Section 5 examines the six principal governance dimensions in turn, presenting each through three lenses: the regulatory obligation that motivates it, a general design principle applicable to any high-risk educational AI system, and its specific implementation in Diotima. Section 6 examines the implications of compliance-by-design for market adoption. Section 7 concludes with the proposition that the regulatory moment in educational technology is also a design opportunity—one that rewards systems built for trust rather than retrofitted to it.

2. AI IN EDUCATION IS HIGH-RISK BY DEFAULT

The EU AI Act's Annex III classification of educational assessment AI as high-risk reflects a straightforward chain of reasoning. Assessment shapes student progression and opportunity. Biased or incorrect AI outputs can materially harm learners. And if automated evaluation is allowed to operate without robust human governance, it progressively displaces the professional judgment that should sit at the centre of any assessment decision.

The high-risk designation applies even where AI is positioned as formative rather than summative. Any system that analyses student responses, structures performance interpretation, and generates feedback can meaningfully influence

outcomes—regardless of how it is labelled. The classification covers not just systems that assign official grades, but any system whose outputs interact with the assessment process in a way that could affect how a learner’s performance is perceived or recorded (EUROPEAN PARLIAMENT, 2024, Annex III, s. 3(a)).

Diotima therefore adopts a conservative, regulator-aligned position from the outset: it is treated as a high-risk system by design. That choice is not defensive—it is strategic. Every architectural decision, every governance mechanism, every operational workflow is oriented by that classification.

Three specific challenges follow from this classification that any serious educational AI platform must address. The first is *hallucination and curriculum drift*: AI language models can generate plausible but factually incorrect or pedagogically misaligned outputs that are not inherently grounded in the curriculum being assessed. In formative assessment contexts, where AI-generated content directly structures how student performance is interpreted, this is not merely a quality problem but a safety one with material consequences for learner outcomes. The second is *automation bias and professional displacement*: well-documented tendencies in human-computer interaction lead users to accept or defer to automated outputs with insufficient critical scrutiny. In educational assessment, this risk takes the specific form of teachers progressively deferring to AI-generated evaluations rather than exercising independent professional judgment—an outcome that Annex III’s oversight requirements are expressly designed to prevent. The third is *the governance architecture deficit*: as documented by FERRARIO ET AL. (2025) and the OECD (2023), the majority of AI products currently deployed in educational settings were not designed for regulatory alignment. Compliance has been treated as a retrospective documentation exercise, creating a structural gap between the obligations that high-risk systems carry and the architectures that purport to meet them.

3. DIOTIMA’S CORE DESIGN PROPOSITION

The most important thing to establish is what Diotima is *not*. It is not an AI grader. It does not autonomously assign marks, determine outcomes, or make binding decisions about student performance. Diotima is a teacher-centred decision-support system for formative assessment—its function is to expand the scope and consistency of formative feedback while preserving the judgment calls that professional expertise requires, not to replace the professional authority at the heart of good teaching.

Each of the three challenges identified in Section 2 has a direct architectural response. Hallucination and curriculum drift are addressed through mandatory grounding of all AI-generated content in approved, licensed source materials—the model cannot generate content that is not explicitly linked to curriculum topics and approved texts. Automation bias and professional displacement are addressed by making human oversight a structural requirement of the system rather than a policy aspiration—no AI output is ever presented to a student without teacher confirmation, and the system is technically incapable of routing unreviewed outputs to learners. The governance architecture deficit is addressed by treating Annex III obligations as the design specification from first principles, so

that compliance evidence is generated as a byproduct of normal operation rather than assembled retrospectively.

In practice, Diotima supports teachers to: generate curriculum-aligned questions grounded in approved source materials; apply consistent, transparent rubrics; provide formative feedback at a scale that would otherwise be impossible; and support student reflection and iterative improvement rather than delivering a single terminal judgment.

The constraint that structures everything else: no AI output is final without human confirmation. Teachers approve all assessment content. Teachers validate all evaluative outputs. The AI proposes; the teacher decides.

4. SYSTEM OVERVIEW

Diotima supports the complete formative assessment lifecycle from initial design through to student reflection and audit. Every stage is engineered to support explainability, traceability, and human oversight—not as retrospective additions, but as structural properties of the workflow itself.

The lifecycle proceeds through seven stages: (1) curriculum-grounded assessment design, in which teachers select topics, subtopics, and cognitive levels and the system generates candidate content from approved source materials; (2) teacher pre-approval of every generated question, expected answer, and rubric before any student sees it; (3) student response and optional AI pre-check, a student-facing feature that provides provisional analysis to support the learner before final submission; (4) conditional rubric visibility, a separate stage in which students receive performance-band information after final submission, designed to support self-assessment without enabling gaming; (5) teacher review, override, and finalisation, during which all AI outputs remain provisional—teachers may accept, edit, augment, or entirely replace any output before it becomes authoritative; (6) student feedback and reflection, in which final teacher-approved feedback is delivered and students are prompted to engage with their learning; and (7) comprehensive logging and auditability, in which every action, inference, and decision is recorded to support traceability, incident response, and post-market monitoring.

Figure 1 illustrates this lifecycle as a continuous flow from assessment design through to audit. These lifecycle stages do not each engage a single governance dimension in isolation. Table 1 maps the six principal governance dimensions to the lifecycle stages at which each is most directly operative. The following sections examine each dimension in turn.

5. GOVERNANCE DIMENSIONS IN PRACTICE

The following sections examine the six principal governance dimensions that structure Diotima’s compliance posture. Each is presented through three lenses: the *regulatory obligation* that motivates it, a *general design principle* applicable to any high-risk educational AI system, and the *specific implementation in Diotima* that demonstrates the principle in practice. While Diotima serves as the case study throughout, the design principles described here are not specific to any single platform—they represent a transferable framework for compliance-by-design in educational AI.




Figure 1: The formative assessment lifecycle. Seven stages from curriculum-grounded assessment design through to comprehensive logging and audit. Stages 1–4 proceed left to right; stages 5–7 return right to left. Every AI-generated output must pass through a teacher approval or review gate (stages 2 and 5) before reaching students.

Table 1: Governance dimensions and lifecycle stage coverage. CG = Curriculum Grounding; HO = Human Oversight; EX = Explainability; MG = Model Governance; DG = Data Governance; PM = Post-Market Monitoring. ✓ indicates primary operative engagement.

Lifecycle Stage	CG	HO	EX	MG	DG	PM
1. Assessment design	✓			✓	✓	
2. Teacher pre-approval	✓	✓	✓			
3. Student response / AI check				✓	✓	
4. Conditional rubric visibility			✓			
5. Teacher review and finalisation		✓	✓			
6. Student feedback and reflection			✓		✓	
7. Logging and auditability				✓	✓	✓

Curriculum Grounding and Risk Scoping

Regulatory obligation. Article 9 of the EU AI Act requires high-risk AI systems to operate within the limits of their intended purpose and to implement risk management measures that address foreseeable misuse and failure modes. For an educational AI system, the most direct foreseeable failure mode is the generation of content that drifts from the curriculum—factually incorrect, pedagogically misaligned, or ungrounded in the materials that the assessment is intended to evaluate.

General principle. High-risk educational AI systems should constrain generative outputs to approved, verified source materials through retrieval-augmented or otherwise grounded generation approaches. Content should be anchored not only to curriculum topics but also to appropriate complexity levels for the target learner population. The high-risk perimeter should be scoped deliberately to the minimum necessary components, so that governance obligations remain tractable and risk propagation is contained.

In Diotima. Diotima addresses hallucination architecturally rather than reactively. Every AI-generated question, answer, and rubric is explicitly linked to curriculum topics and subtopics, specified learning outcomes, cognitive levels aligned to the target learner population, and approved, licensed source-of-truth materials.

This grounding serves three functions simultaneously. It prevents the model from generating content that drifts from

the curriculum. It ensures that every item can be traced back to its source, giving teachers the visibility they need to make informed approval decisions. And it provides a clear basis for explaining why an item exists and what it is assessing—which is exactly what regulators and institutions require from a high-risk system.

A second architectural decision is the deliberate scoping of the high-risk AI perimeter within the system. Diotima as a whole operates under Annex III; within it, the high-risk AI inference is concentrated in two components—the Question, Answer and Rubric Generation Engine, and the AI Inference and Feedback Engine—while supporting systems (user interfaces, workflow orchestration, data storage, and integration components) do not themselves execute high-risk AI functionality and are governed accordingly. Narrowing the internal high-risk footprint in this way reduces risk propagation, simplifies governance, and makes compliance obligations clearer and more tractable.

Lifecycle coverage. Grounding operates primarily at stages 1 and 2: content is anchored to curriculum materials at the point of generation and verified against those materials during teacher pre-approval.

Human Oversight

Regulatory obligation. Article 14 of the EU AI Act requires high-risk AI systems to be designed and developed so that they can be effectively overseen by natural persons (i.e.

human professionals) during their use. The obligation extends to enabling oversight persons to understand the system’s capabilities and limitations, to detect and address failures, and to intervene on, override, or halt the system’s operation.

General principle. Human oversight should be enforced as a structural property of the system workflow, not declared as a policy layer applied after the fact. The system should be technically incapable of routing AI-generated outputs to affected persons—in educational contexts, students—without prior review and confirmation by a qualified human professional. This ensures that accountability remains with the human decision-maker rather than being diffused across an automated pipeline.

In Diotima. Teachers must approve every generated assessment item before it reaches students. AI rubric placements are always provisional and are never presented as final. Teachers can accept, edit, override, or entirely replace any AI output at any stage, and only teacher-confirmed feedback is stored as the authoritative record. Students never receive AI-generated feedback that a teacher has not reviewed. That is not a policy aspiration—it is a technical constraint embedded in the system.

The consequences of this design are significant. Professional autonomy is preserved. The structural risk of over-reliance on AI outputs is prevented rather than mitigated after the fact. And accountability stays where it belongs—with the human professional, not the algorithm.

Lifecycle coverage. Human oversight is operative primarily at stages 2 and 5: the pre-approval checkpoint before content reaches students, and the review and finalisation stage before feedback is delivered.

Role-Differentiated Explainability

Regulatory obligation. Article 13 of the EU AI Act requires high-risk AI systems to be sufficiently transparent that users can interpret the system’s outputs and use them appropriately. Transparency obligations extend both to the users operating the system—in this case, teachers—and, where relevant, to the persons affected by its outputs—in this case, students.

General principle. Explainability should be differentiated by stakeholder role. System operators—typically teachers or institutional administrators—require full transparency into the reasoning behind AI outputs, including source materials, evaluation criteria, and the basis for any automated judgment. Affected persons—typically students—require contextually appropriate transparency that supports their engagement with the process without undermining the integrity of the assessment itself. A single undifferentiated transparency model risks either overwhelming students with technical detail or depriving professionals of the information they need to exercise meaningful oversight.

In Diotima. For teachers, the system provides full visibility into the rubric and the source text from which assessment content was derived, so that every output can be understood in terms of the curriculum materials that grounded it. The pattern of rejection reasons and overrides constitutes an ongoing signal about model performance—a practical feedback loop that informs model governance decisions (see Section 5.6).

For students, the explainability design is more deliberate.

Rubric information is revealed conditionally and progressively: students see performance bands after a genuine attempt, not upfront. Full criteria are not disclosed in advance. This supports transparency about how responses are being interpreted while preserving the pedagogical integrity of the assessment, and it orients feedback toward improvement and next steps rather than terminal judgment.

Lifecycle coverage. Explainability operates wherever AI outputs meet a human: teachers at stages 2 and 5, students at stages 4 and 6.

Model Governance

Regulatory obligation. Article 9 of the EU AI Act requires high-risk AI providers to implement risk management measures throughout the system lifecycle, including identification and analysis of known and foreseeable risks, and estimation and evaluation of risks that may emerge from actual use. For systems using third-party foundation models, this obligation extends to the selection, evaluation, and ongoing monitoring of those models.

General principle. For systems built on third-party foundation models, model selection should be treated as a governance decision, not a technical one. A structured benchmark suite—covering knowledge and reasoning, factual accuracy, instruction adherence, bias, and safety—provides the evidential basis for initial selection and ongoing monitoring. Model versions should be logged with every inference to support audit, incident investigation, and drift detection. No student data should be used for model training or fine-tuning.

In Diotima. Diotima uses third-party pretrained foundation models rather than training its own. That is a deliberate and defensible choice—but it does not mean model selection is a technical decision. In Diotima, model choice is treated as a governance decision.

Models are evaluated and selected against a structured benchmark suite. Knowledge and reasoning capabilities are assessed using the Massive Multitask Language Understanding benchmark (MMLU; HENDRYCKS ET AL., 2021)—a 57-subject evaluation spanning STEM, humanities, and professional domains—and the Graduate-Level Google-Proof Q&A benchmark (GPQA; REIN ET AL., 2023), which tests expert-level reasoning across biology, chemistry, and physics. Factual accuracy and hallucination propensity are evaluated using LongFact (WEI ET AL., 2024), which probes the reliability of long-form factual claims across a wide range of topics. Instruction-following behaviour is assessed using the Instruction Following Evaluation benchmark (IFEval; ZHOU ET AL., 2023), a verifiable evaluation of adherence to structured natural language instructions. Bias and fairness are assessed using the Bias Benchmark for Question Answering (BBQ; PARRISH ET AL., 2022), a hand-built benchmark targeting question-answering behaviour across nine social dimensions including age, gender, and race. Toxicity and safety are assessed using RealToxicityPrompts (GEHMAN ET AL., 2020), which evaluates the tendency of models to produce harmful content across a corpus of naturally occurring prompts, and SimpleSafetyTests (VIDGEN ET AL., 2023), a test suite of prompts spanning five harm areas designed to identify critical safety risks such as

unsafe advice and toxic content generation.

No student data is used for model training at any stage. Model versions are logged with every inference, creating the audit trail needed for incident investigation, drift detection, and post-market monitoring. This approach enables evidence-based model selection, ongoing comparison as new models emerge, and safe substitution of model versions without architectural change.

Lifecycle coverage. Model governance is operative at stages 1, 3, and 7: model selection determines what is generated at the design stage; model inference is used during the student response and AI pre-check stage; and model version logging at stage 7 supports post-market monitoring.

Data Governance

Regulatory obligation. Article 10 of the EU AI Act requires high-risk AI systems to be trained, validated, and tested on data that meets quality criteria including relevance, representativeness, and freedom from errors and bias. More broadly, data governance for systems processing personal data must be compatible with applicable data protection law, including the General Data Protection Regulation (GDPR). For educational AI systems processing student data—a category that warrants particular care given the vulnerability of the population—data minimisation and purpose limitation are foundational requirements.

General principle. Data minimisation and purpose limitation should be foundational design constraints, not compliance additions applied after the fact. Educational AI systems should process only the data strictly necessary for their stated purpose, with no repurposing of student data for model training, behavioural profiling, or secondary commercial use. Where third-party educational content is used as source material, its governance—including storage, processing, and derivative use—should be treated as a first-class compliance obligation under explicit licensing agreements.

In Diotima. The system processes only what is necessary: student responses, assessment configuration, and the metadata required to deliver formative feedback and support audit. There is no collection of sensitive personal categories, no behavioural profiling, and no use of student data to fine-tune or retrain the underlying models.

Retention limits are defined and aligned with institutional Data Protection Impact Assessment (DPIA) requirements. Educational content—curriculum documents, licensed textbooks, and approved source materials—is managed under explicit publisher agreements covering storage, processing, and derivative use within the scope of Diotima’s purpose.

The effect is a data governance posture that protects every stakeholder simultaneously: students, whose assessment data is not repurposed; teachers, whose professional actions are logged for audit but not exposed inappropriately; publishers, whose content is used under clear and documented licence; and institutions, whose DPIA obligations are supported rather than complicated by the platform.

Lifecycle coverage. Data governance is operative at stages 1, 3, 6, and 7: at the point of accessing source ma-

terials, when processing student responses, when delivering and storing feedback, and at the comprehensive logging stage.

Post-Market Monitoring

Regulatory obligation. Article 72 of the EU AI Act requires providers of high-risk AI systems to establish and document a post-market monitoring system that actively collects and analyses data on the system’s performance throughout its operational lifetime. Post-market monitoring is a legal obligation, not a discretionary quality measure—and it is where many AI governance frameworks, in practice, fail to deliver on their stated commitments.

General principle. Post-market monitoring should be designed so that evidence of safe operation is generated as a byproduct of normal system use, rather than requiring separate data collection or retrospective assembly. In educational AI systems, structured feedback from professional users—teacher rejection reasons, override patterns, and incident reports—provides a particularly rich source of real-world performance data, capturing information about model behaviour in authentic pedagogical contexts that synthetic benchmarks alone cannot replicate.

In Diotima. Diotima embeds post-market monitoring from day one as an operational discipline rather than an afterthought.

Models are regularly re-evaluated against the benchmark suite to detect drift and performance degradation. Anomaly detection flags distribution shifts in AI outputs before they reach users. Teachers have structured channels to report harmful or incorrect outputs, with time-bound remediation workflows. The system is designed for registration in the EU high-risk AI database, with version update processes in place.

The teacher approval and rejection workflow provides a continuous signal about model performance in real educational contexts. Rejection reasons and override patterns are reviewed by human governance personnel and manually entered into the risk register and model governance process—a structured feedback loop grounded in actual pedagogical use rather than synthetic evaluation. The system is designed so that this evidence of safe operation is generated as a byproduct of normal workflow rather than requiring separate data collection.

Lifecycle coverage. Post-market monitoring is operative principally at stage 7, where comprehensive logging creates the data infrastructure for ongoing monitoring, but it is also informed continuously by teacher actions at stages 2 and 5.

6. COMPLIANCE-BY-DESIGN AND THE PATH TO ADOPTION

The EU AI Act’s entry into force marks a structural shift in the conditions for market participation in educational technology. The obligations imposed by Annex III—relating to human oversight, transparency, data governance, model evaluation, and post-market monitoring—are not documentation requirements. They are design requirements. Products that treat them as the former face a growing and ultimately unbridgeable gap between declared and demonstrable compliance.

This gap is not hypothetical. Analysis of the educational

AI supplier market suggests that very few products currently deployed in formal educational settings would meet the substantive requirements of the Act as currently written—and this pattern is not confined to the EU. The UK Department for Education published its *Generative AI: Product Safety Standards* (DEPARTMENT FOR EDUCATION, 2025) setting out the capabilities and features that AI products must meet to be considered safe for use in educational settings. Industry analysis by the British Educational Suppliers Association indicates that as of early 2026, no supplier has been confirmed as fully compliant with these standards, including products that have received government endorsement (BESA, personal communication, 2026). The structural reason is the same in both markets: compliance was not designed in.

Diotima’s architecture responds to this directly. Because Annex III obligations were treated as the design specification from the outset, the compliance posture is not asserted—it is verifiable. Each governance dimension documented in this paper corresponds to a concrete architectural or operational feature that can be examined, audited, and validated by an institution, a regulator, or a conformity assessment body. This is what distinguishes compliance-by-design from post-hoc compliance: the former generates evidence as a byproduct of normal operation; the latter requires it to be produced retrospectively, from a system not designed to produce it.

The institutional implications are direct. Educational institutions considering AI-assisted formative assessment do not need to choose between regulatory safety and pedagogical capability. A system designed from first principles for Annex III compliance can be simultaneously safe, transparent, and educationally serious. The architecture described in this paper demonstrates that compliance, far from constraining what the technology can do, is the foundation on which trustworthy educational AI must be built.

7. CONCLUSION

The central proposition of this paper is that the EU AI Act’s high-risk classification of educational assessment AI, rather than constraining what can be built, defines what must be built. A system that takes Annex III obligations seriously as design requirements—grounding its outputs in approved curriculum materials, enforcing human oversight structurally, differentiating explainability by role, governing model selection and data use to evidential standards, and embedding post-market monitoring as an operational discipline—is not a constrained system. It is a more trustworthy, more auditable, and more professionally empowering one.

Diotima represents one implementation of this proposition. The architecture described in this paper demonstrates that compliance-by-design is not a theoretical position but a practically achievable design outcome—one that preserves teacher professional authority, supports pedagogically serious formative assessment, and generates the evidence of safe operation that institutions and regulators require.

Architectural foundations, however, do not on their own answer every operational question. Concrete metrics for measuring compliance in practice—hallucination rates, validation

that teacher review effort is substantive rather than perfunctory, prompt governance and audit trails, and closed-loop mechanisms that feed teacher corrections back into model updates—are substantive issues in their own right. A companion paper will address these operational controls and the measurement frameworks required to evidence compliance at scale.

The regulatory moment in educational technology creates an obligation that is also an opportunity: to build AI for education that earns institutional trust rather than requests it, and that places the professional judgment of teachers at the centre of every consequential decision about learners.

REFERENCES

- BLOOM B. S., ENGELHART M. D., FURST E. J., HILL W. H. and KRATHWOHL D. R. (1956) *Taxonomy of Educational Objectives: The Classification of Educational Goals. Handbook I: Cognitive Domain*. David McKay, New York.
- DEPARTMENT FOR EDUCATION (2025) *Generative AI: Product Safety Standards*, GOV.UK. Available at: <https://www.gov.uk/government/publications/generative-ai-product-safety-standards> (updated 19 January 2026).
- EUROPEAN PARLIAMENT AND COUNCIL OF THE EUROPEAN UNION (2024) Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act), *Official Journal of the European Union*, L 2024/1689.
- FERRARIO A., TULLI S., BRACCO F. and ZANATTO D. (2025) Testing the applicability of a governance checklist for high-risk AI-based learning outcome assessment in Italian universities under the EU AI Act Annex III, *Frontiers in Artificial Intelligence*, 8, 1718613.
- GEHMAN S., GURURANGAN S., SAP M., CHOI Y. and SMITH N. A. (2020) RealToxicityPrompts: Evaluating Neural Toxic Degeneration in Language Models, *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 3356–3369.
- HENDRYCKS D., BURNS C., BASART S., ZOU A., MAZEIKA M., SONG D. and STEINHARDT J. (2021) Measuring Massive Multi-task Language Understanding, *International Conference on Learning Representations (ICLR 2021)*, arXiv:2009.03300.
- NOVAL CONSULTING (2025) *Diotima AI Assessment Platform: Annex III High-Risk AI Compliance Dossier, v1.1-draft*. Internal technical documentation, Noval Consulting, Dublin.
- OECD (2023) *Emerging governance of generative AI in education*, in *OECD Digital Education Outlook 2023*, OECD Publishing, Paris.
- PARRISH A., CHEN A., NANGIA N., PADMAKUMAR V., PHANG J., THOMPSON J., HTUT P. M. and BOWMAN S. (2022) BBQ: A hand-built bias benchmark for question answering, *Findings of the Association for Computational Linguistics: ACL 2022*, pp. 2982–2996.
- REIN D., HOU B. L., STICKLAND A. C., PETTY R., PANG R. Y., DIRANI N., MICHAEL J. and BOWMAN S. R. (2023) GPQA: A Graduate-Level Google-Proof Q&A Benchmark, *arXiv preprint arXiv:2311.12022*.
- VIDGEN B., THRUSH T., WASEEM Z. and KIELA D. (2023) SimpleSafetyTests: a Test Suite for Identifying Critical Safety Risks in Large Language Models, *arXiv preprint arXiv:2311.08370*.

WEI J., YANG C., SONG Y., LU P., HU E., TRAN D., PENG D., LIU R., HUANG D., DU Y. and LE Q. V. (2024) Long-form factuality in large language models, *arXiv preprint* arXiv:2403.18802.

ZHOU J., LU P., MISHRA S., BRAHMA S., BASU S., LUAN Y., ZHOU D. and HOU L. (2023) Instruction-Following Evaluation for Large Language Models, *arXiv preprint* arXiv:2311.07911.